

```
package main

import "fmt"

func main() {
    fmt.Println("Are you testing your Observability?")
    fmt.Println("    --- Metrics Edition ---    ")

    fmt.Println("    GoDays 22.01.2020, Berlin    ")
}
```



@bwplotka  
@kakkoyun



# Thanos

@bwplotka  
@kakkoyun

HTTP  
/lb

httputil.ReverseProxy.ServerHTTP(...)

lbtransport.RoundTripper

Discoverer.Targets()

RoundRobinPicker.Pick()

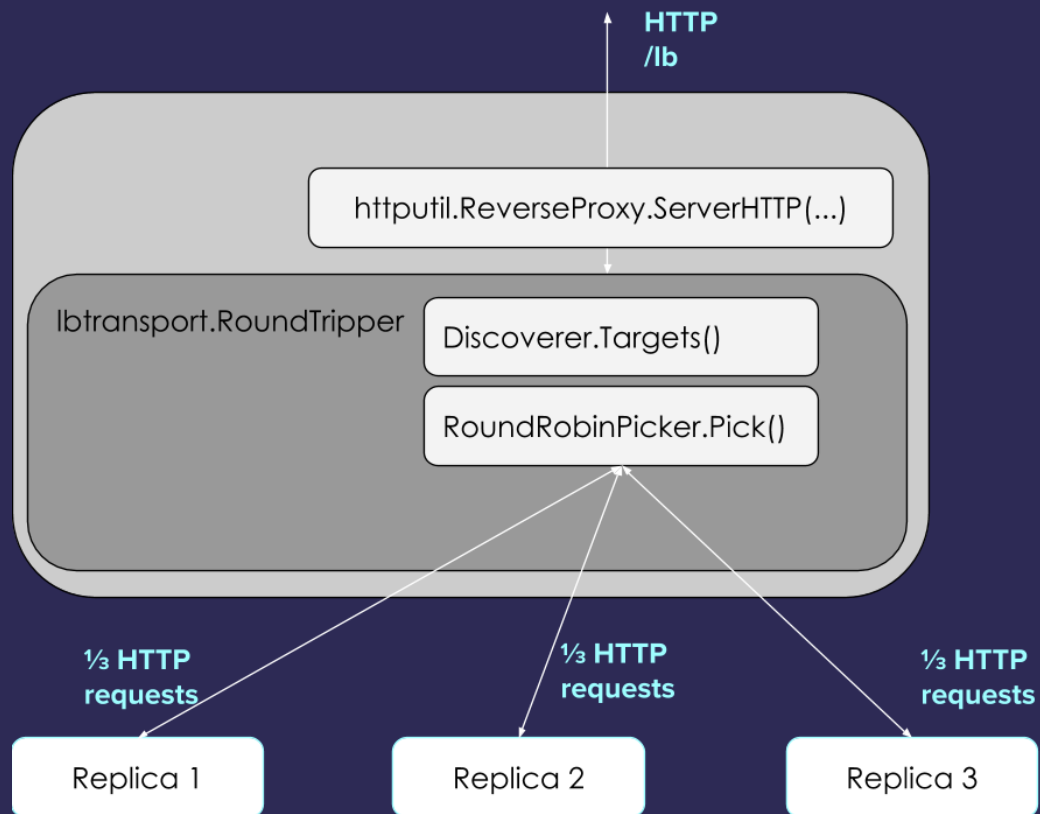
1/3 HTTP  
requests

1/3 HTTP  
requests

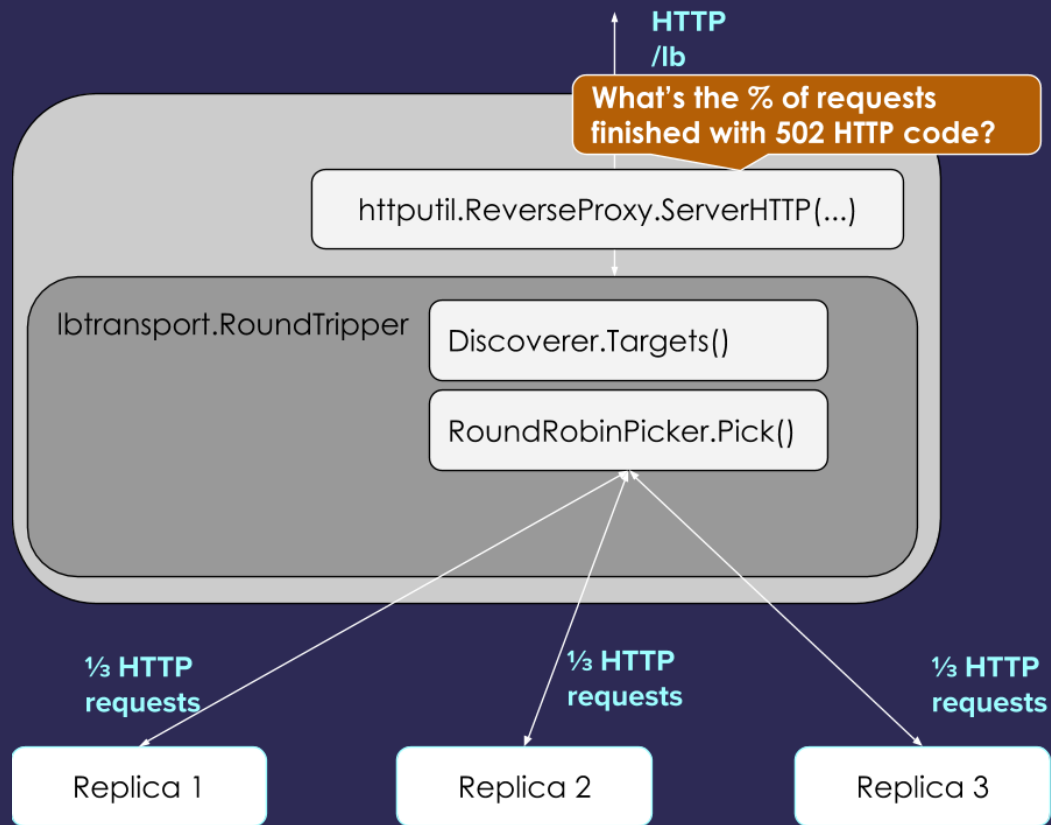
1/3 HTTP  
requests

@hwplotka  
@kakkoyun

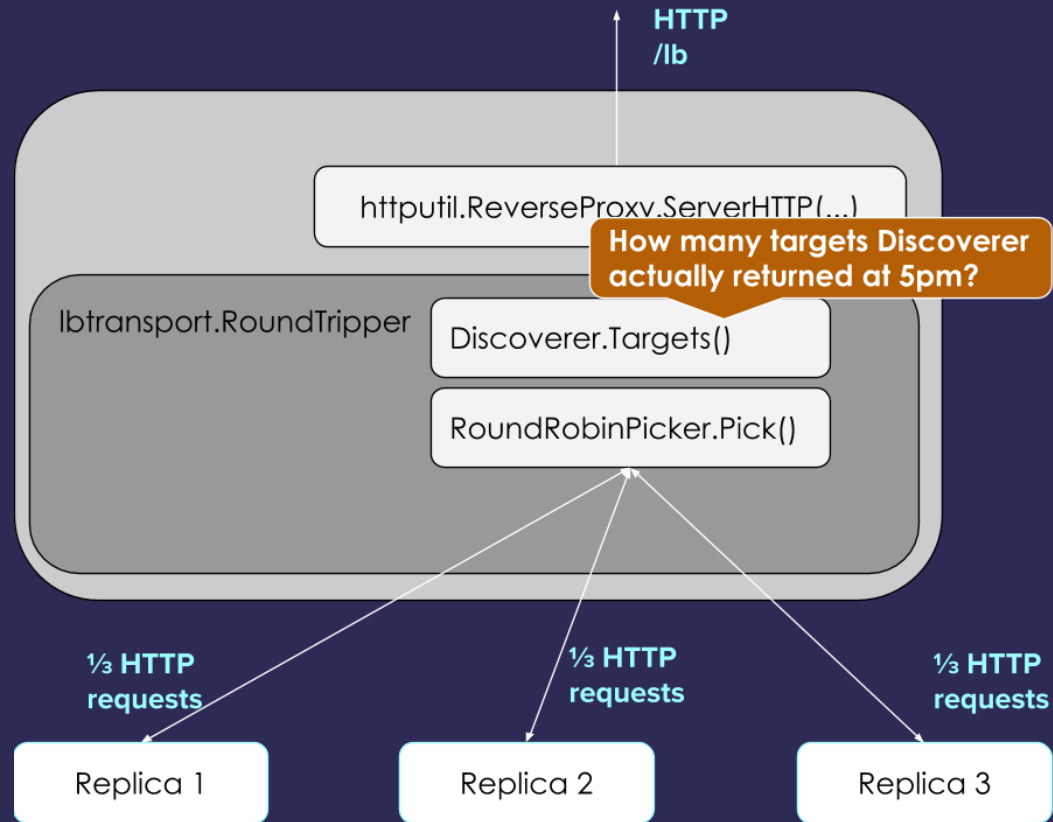
## Wait... Anything missing?



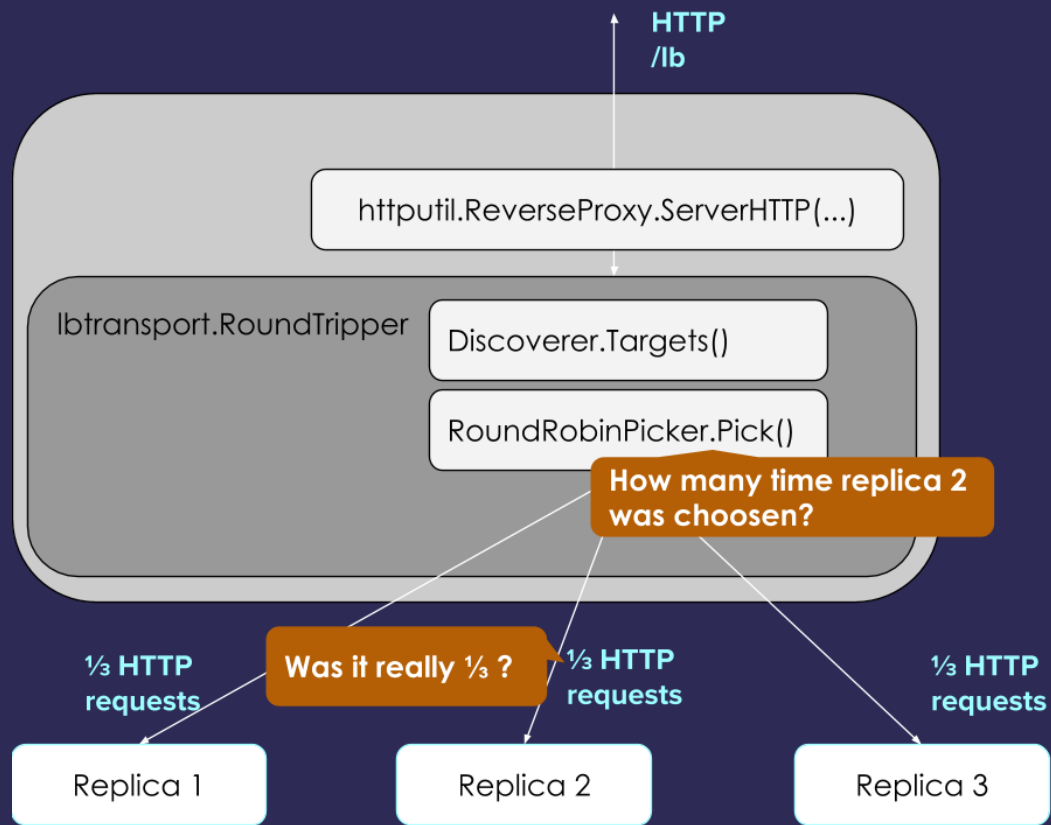
# Wait... Anything missing?



# Wait... Anything missing?



# Wait... Anything missing?





/lb

What latency LB introduces  
for most requests?

httputil.ReverseProxy.ServerHTTP(...)

lbtransport.RoundTripper

Discoverer.Targets()

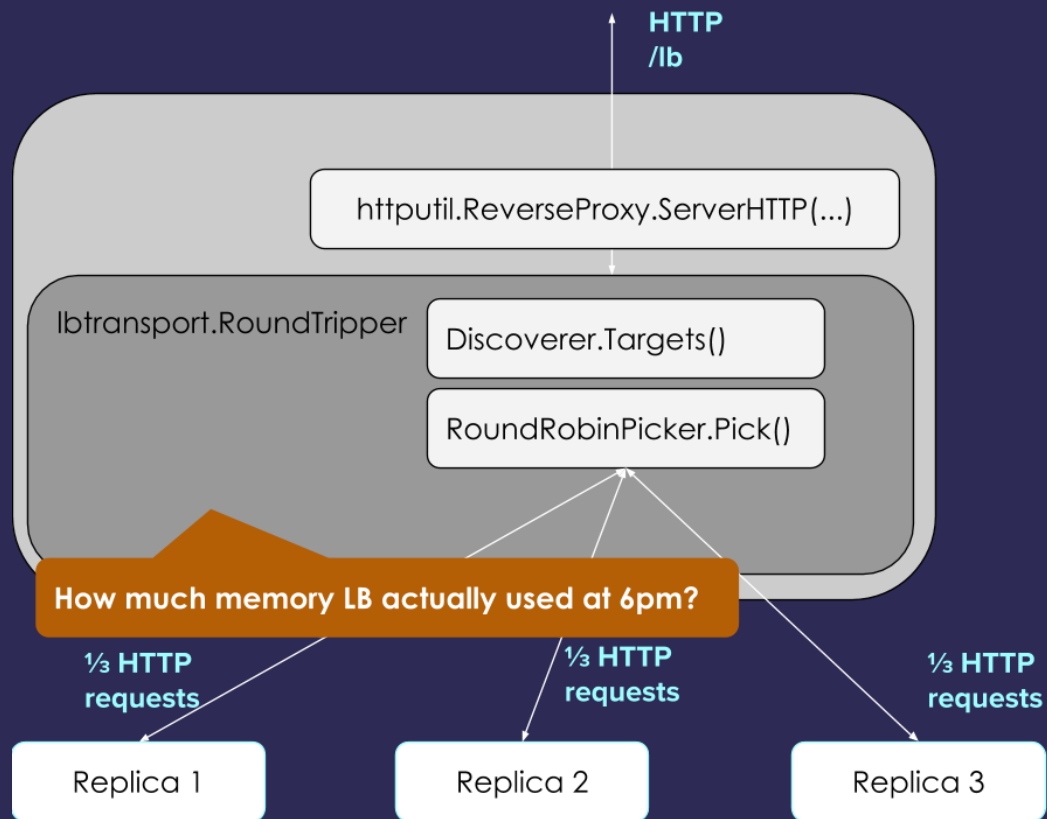
RoundRobinPicker.Pick()

Or maybe it's replica 1 that is slow?

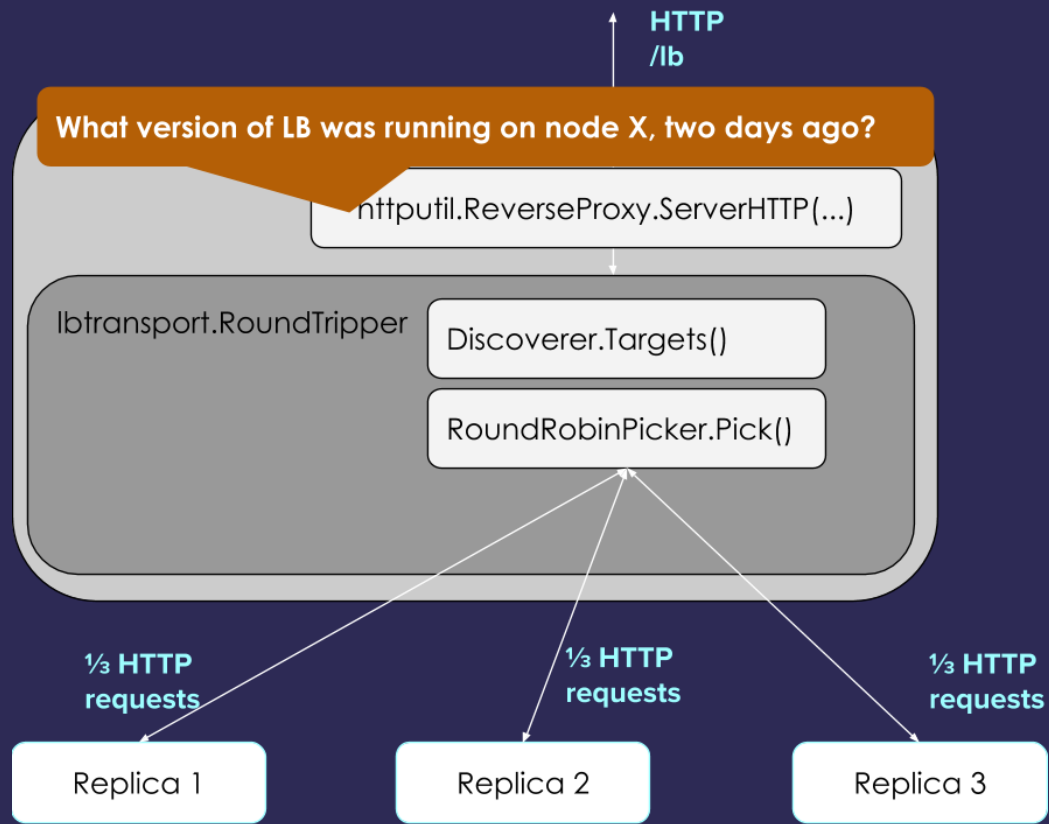
1% HTTP  
requests

1/3 HTTP  
requests  
@bwplotka  
@kakkoyun

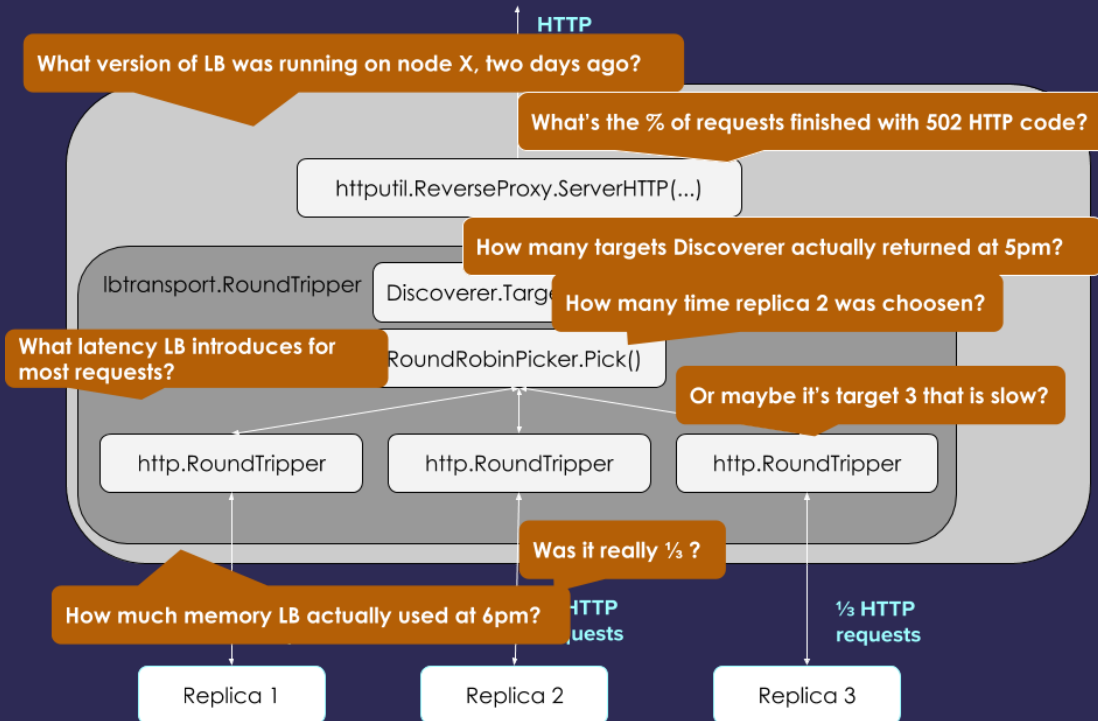
# Wait... Anything missing?



# Wait... Anything missing?



# We need Monitoring!



# Let's instrument our LB with Prometheus metrics!

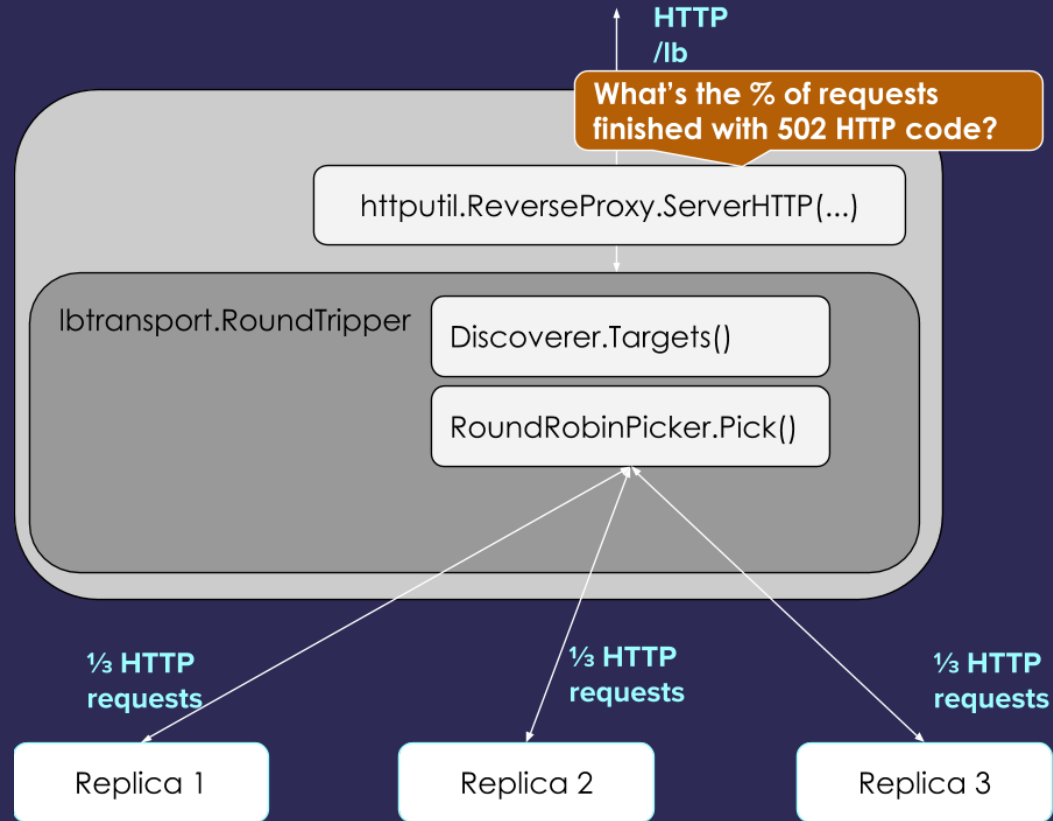
- Cheap
- Near Real Time
- Actionable (Alert-able)



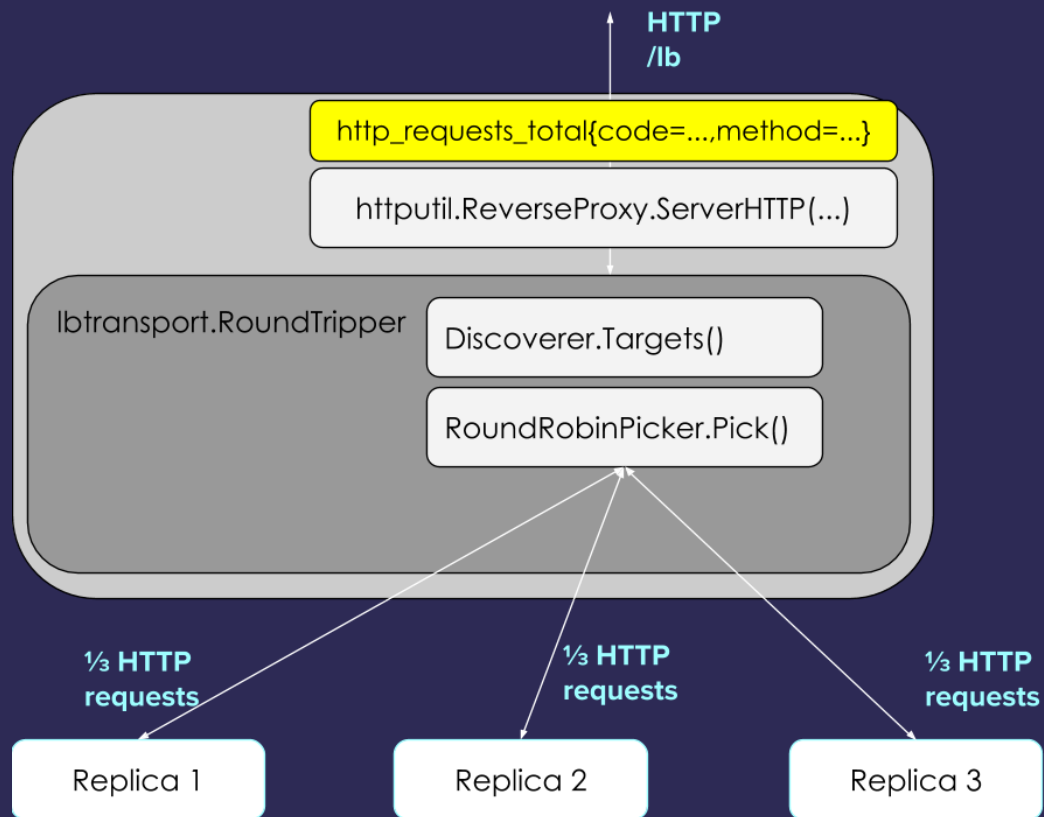
<http://prometheus.io>

@bwplotka  
@kakkoyun

# Let's instrument our LB with Prometheus metrics!



# Let's instrument our LB with Prometheus metrics!



## Server HTTP request counter

```
// Increment our counter with written status code and request method.
serverRequestsTotal.WithLabelValues(statusRec.Status(), r.Method).Add(1)
}

func init() {
    prometheus.MustRegister(serverRequestsTotal)
}

mux := http.NewServeMux()
mux.Handle("/metrics", promhttp.Handler())
mux.Handle("/lb", ...)

// Other handlers...

srv := &http.Server{Handler: mux}

// Run server...
```

Add /metrics handler.

@bwplotka  
@kakkoyun



## From code to graph

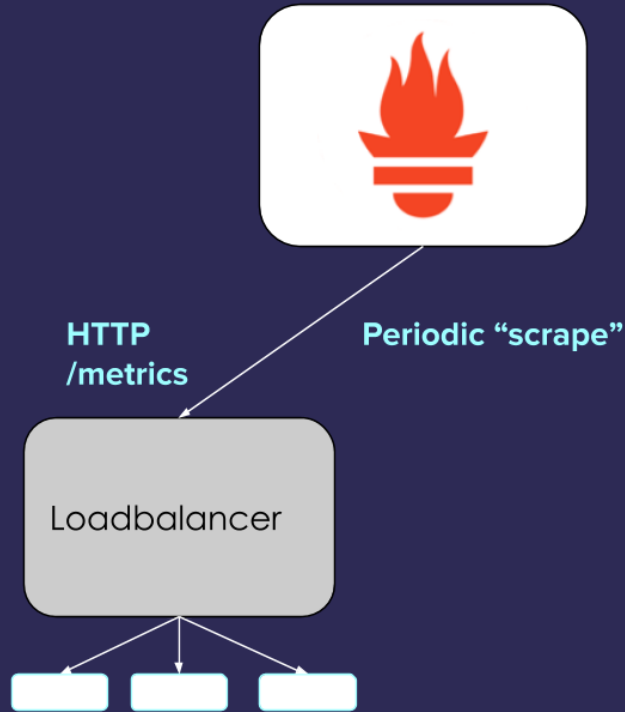
```
var (  
    serverRequestsTotal = prometheus.NewCounterVec(  
        prometheus.CounterOpts{  
            Name: "http_requests_total",  
            Help: "Tracks the number of HTTP requests.",  
        }, []string{"code", "method"},  
    )  
  
    // Part of response from /metrics HTTP endpoint.  
    # HELP http_requests_total Tracks the number of HTTP requests.  
    # TYPE http_requests_total counter  
    http_requests_total{code="200", method="get"} 1089  
    http_requests_total{code="500", method="get"} 46
```

It is now exposed under /metrics.

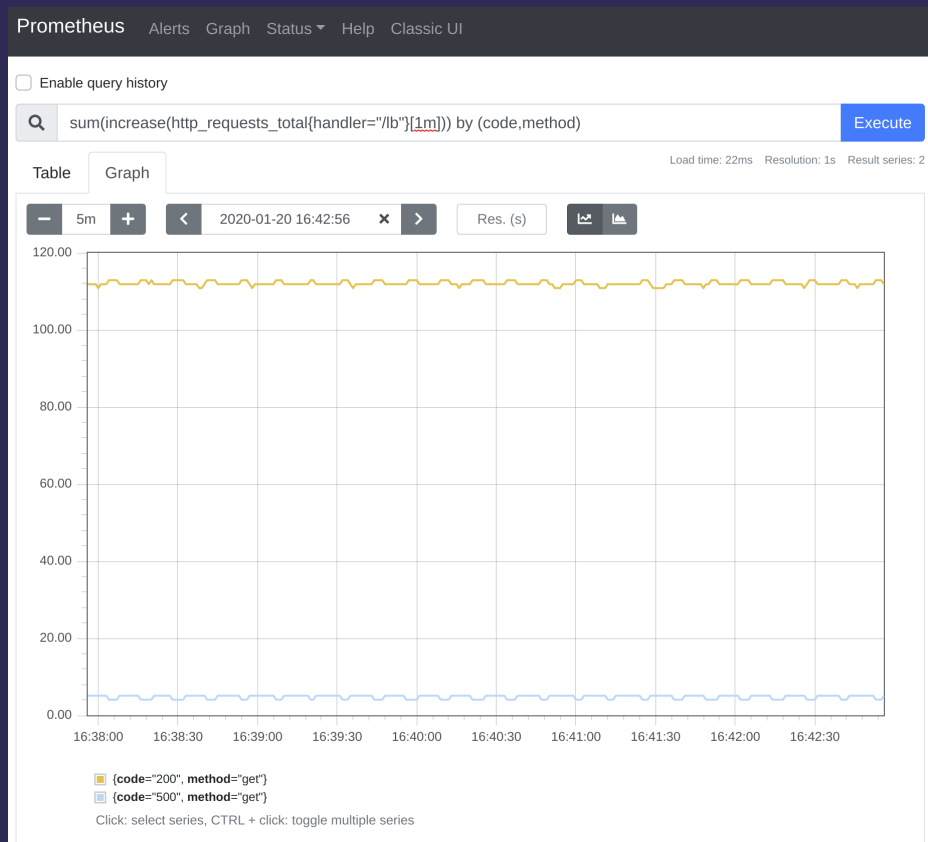
@bwplotka  
@kakkoyun

# Prometheus can now collect metrics from our Loadbalancer

*Pulling e.g every 15 seconds*



# Graph: Prometheus UI





+



Pitfalls



@bwplotka  
@kakkoyun

# Pitfall #1: Global Registry

"magic is bad; global state is magic" by Peter Bourgon



@bwplotka  
@kakkoyun

## Pitfall #1: Getting rid of globals.

```
)  
metrics3 := NewServerMetrics(  
    prometheus.WrapWithLabels(prometheus.Labels{"handler":"/three"}, reg),  
)  
  
// For endpoint /one:  
metrics1.requestsTotal.WithLabelValues(statusRec.Status(), r.Method).Add(1)  
  
// For endpoint /two:  
metrics2.requestsTotal.WithLabelValues(statusRec.Status(), r.Method).Add(1)  
  
// For endpoint /three:  
metrics3.requestsTotal.WithLabelValues(statusRec.Status(), r.Method).Add(1)  
  
# HELP http_requests_total Tracks the number of HTTP requests.  
# TYPE http_requests_total counter  
http_requests_total{handler="/one", code="200", method="get"} 1445  
http_requests_total{handler="/one", code="500", method="get"} 23  
http_requests_total{handler="/two", code="200", method="get"} 445  
http_requests_total{handler="/two", code="500", method="get"} 0  
http_requests_total{handler="/three", code="200", method="get"} 645  
http_requests_total{handler="/three", code="500", method="get"} 40
```

We can have request counter per handler (:

@bwplotka  
@kakkoyun

## Pitfall #2: No Tests For Metrics



## Pitfall #2: No testing: Let's add tests!

```
// other cases...
} {
    if ok := t.Run("", func(t *testing.T) {
        // Prepare mocks...

        rec := httptest.NewRecorder()
        lb.ServeHTTP(
            rec,
            httptest.NewRequest("GET", "http://mocked", nil),
        )
        testutil.Equals(t, tcase.expectedCode, rec.Code)
        testutil.Equals(t, tcase.expectedHost, rec.Header().Get("X-lb-host"))

        // == RUN    TestLoadbalancer/#00
        //      --- FAIL: TestLoadbalancer/#00 (0.00s)
        //      transport_test.go:190:
        //
        //          exp: 1
        //
        //          got: 0
        //      FAIL
        //
        testutil.Equals(
            t,
            tcase.requestsCode502,
            promtestutil.ToFloat64(metrics.requestsTotal.WithLabelValues("502", "get")),
        )
        testutil.Equals(t, 1, promtestutil.CollectAndCount(lb.metrics.requestsTotal))

    }); !ok {
        return
    }
}
```

Now, test will catch our bug!

@bwplotka  
@kakkoyun



## Pitfall #3: Lack of Consistency

The Four Golden Signals, USE method, RED method etc...

- **R**: Requests per second (saturation).
- **E**: Errors per second.
- **D**: Duration (tail latency).

## Pitfall #3: Consistency

```
type ServerMetrics struct {
    requestsTotal    *prometheus.CounterVec
    requestDuration  *prometheus.HistogramVec
}

// NewServerMetrics provides ServerMetrics.
func NewServerMetrics(reg prometheus.Registerer) *ServerMetrics {
    m := &ServerMetrics{
        requestsTotal: prometheus.NewCounterVec(
            prometheus.CounterOpts{
                Name: "http_requests_total",
                Help: "Tracks the number of HTTP requests.",
            }, []string{"code", "method"}),
        requestDuration: prometheus.NewHistogramVec(
            prometheus.HistogramOpts{
                Name:    "http_request_duration_seconds",
                Help:    "Tracks the latencies for HTTP requests.",
                Buckets: []float64{0.001, 0.01, 0.1, 0.3, 0.6, 1, 3, 6, 9, 20, 30, 60, 90, 120},
            }, []string{"code", "method"}),
    }
    reg.MustRegister(m.requestsTotal, m.requestDuration)
    return m
}

// Top level ServeHTTP handler.
func ServeHTTP(w http.ResponseWriter, r *http.Request) {
    start := time.Now()
    defer metrics.requestDuration.WithLabelValues(statusRec.Status(), r.Method).Observe(time.Since(start))

    statusRec := newStatusRecorder(w)
    next.ServeHTTP(statusRec, r)

    // Increment our counter with written status code and request method.
```

Red method satisfied

@bwplotka  
@kakkoyun

## **Pitfall #4: Naming: Not conforming naming convention**

**There is an official documentation on naming conventions**

<https://prometheus.io/docs/practices/naming/#metric-and-label-naming>

@bwplotka  
@kakkoyun

## Pitfall #4: Naming: Not conforming naming convention

`http_request_duration_seconds`

`node_memory_usage_bytes`

`http_requests_total`

`process_cpu_seconds_total`

`build_info`

`_info` suffix for metadata metrics

@bwplotka  
@kakkoyun

## Pitfall #4: Naming: stability

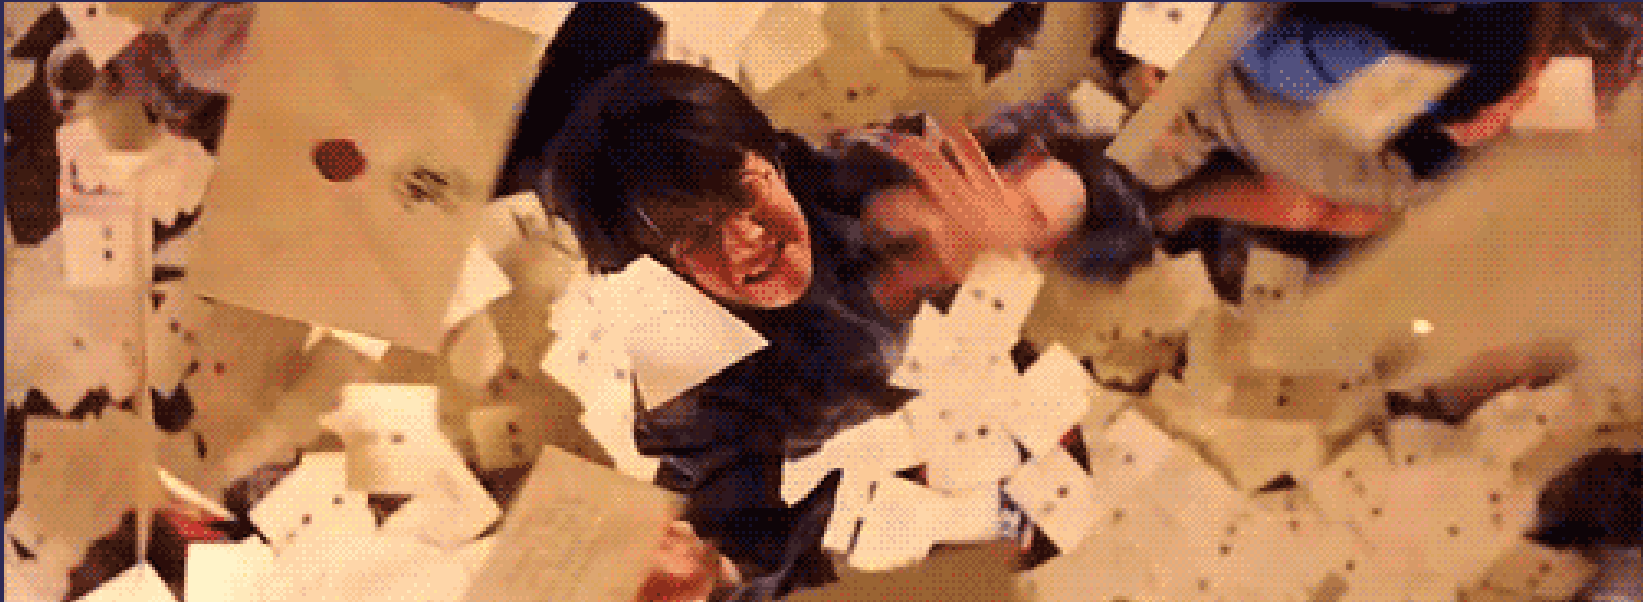
```
// NewServerMetrics provides ServerMetrics.
func NewServerMetrics(reg prometheus.Registerer) *ServerMetrics {
    m := &ServerMetrics{
        requestsTotal: prometheus.NewCounterVec(
            prometheus.CounterOpts{
                Name: "http_protocol_requests_total",
                Help: "Tracks the number of HTTP requests.",
            }, []string{"code", "method"},
        ),
    }
    reg.MustRegister(m.requestsTotal)
    return ins
}

## BOOM!💣 This alert will never fire but also will not fail!
## Rename can cause issues like this in Alerts, Recording rules, Dashboards and more..
alert: HttpToMany502Errors
  expr: |
    sum(rate(http_requests_total{status="502"}[1m])) /
    sum(rate(http_requests_total[1m])) * 100 > 5
  for: 5m
  labels:
    severity: error
  annotations:
    summary: "HTTP errors 502 (instance {{ $labels.instance }})"
    description: |
      "Too many HTTP requests with status 502 (> 5%)\n  VALUE = {{ $value }}\n
      LABELS: {{ $labels }}"
```

Ups...

@bwplotka  
@kakkoyun

## Pitfall #5: Cardinality: Unbounded labels



```
# HELP conntrack_dialer_conn_failed_total Total number of connections failed to dial by the dialer.
# TYPE conntrack_dialer_conn_failed_total counter
conntrack_dialer_conn_failed_total{reason="<nil>"} 1
conntrack_dialer_conn_failed_total{reason="lookup example.com on localhost: err..."} 1
conntrack_dialer_conn_failed_total{reason="lookup thanos.io on 8.8.8.8: err..."} 1
conntrack_dialer_conn_failed_total{reason="lookup redhat.com on localhost: err..."} 1
conntrack_dialer_conn_failed_total{reason="syscall: unimplemented EpollWait"} 1
conntrack_dialer_conn_failed_total{reason="syscall.SIGINT: series of unfortunate things happened"} 1
conntrack_dialer_conn_failed_total{reason="unix: test returned fd in TestEpoll"} 1
conntrack_dialer_conn_failed_total{reason="Invalid value for argument: client: nil"} 1
```

Let's check out our metrics... 🤔

@bwplotka  
@kakkoyun



```

        metrics.ConnFailedReasonLabel.SetValue(dialErrToReason(err)).Inc()
        return conn, err
    }

    return &clientConnTracker{...}, nil
}

func dialErrToReason(err error) string {
    var e *net.OpError
    if errors.As(err, &e) {
        switch nestErr := e.Err.(type) {
        case *net.DNSError:
            return failedResolution
        case *os.SyscallError:
            if nestErr.Err == syscall.ECONNREFUSED {
                return failedConnRefused
            }

            return failedUnknown
        }

        if e.Timeout() {
            return failedTimeout
        }
    } else if err == context.Canceled || err == context.DeadlineExceeded {
        return failedTimeout
    }

    return failedUnknown
}

```

## **Pitfall #6: Cardinality: Histogram Explosion**

## Pitfall #6: Histogram Cardinality Explosion

```
# HELP http_request_duration_seconds Tracks the latencies for HTTP requests.
# TYPE http_request_duration_seconds histogram
http_request_duration_seconds_bucket{code="200",handler="/metrics",method="get",le="0.001"} 0
http_request_duration_seconds_bucket{code="200",handler="/metrics",method="get",le="0.01"} 18
http_request_duration_seconds_bucket{code="200",handler="/metrics",method="get",le="0.1"} 18
http_request_duration_seconds_bucket{code="200",handler="/metrics",method="get",le="0.3"} 18
http_request_duration_seconds_bucket{code="200",handler="/metrics",method="get",le="0.6"} 18
http_request_duration_seconds_bucket{code="200",handler="/metrics",method="get",le="1"} 18
http_request_duration_seconds_bucket{code="200",handler="/metrics",method="get",le="3"} 18
http_request_duration_seconds_bucket{code="200",handler="/metrics",method="get",le="6"} 18
http_request_duration_seconds_bucket{code="200",handler="/metrics",method="get",le="9"} 18
http_request_duration_seconds_bucket{code="200",handler="/metrics",method="get",le="20"} 18
http_request_duration_seconds_bucket{code="200",handler="/metrics",method="get",le="30"} 18
http_request_duration_seconds_bucket{code="200",handler="/metrics",method="get",le="40"} 18
http_request_duration_seconds_bucket{code="200",handler="/metrics",method="get",le="50"} 18
http_request_duration_seconds_bucket{code="200",handler="/metrics",method="get",le="60"} 18
http_request_duration_seconds_bucket{code="200",handler="/metrics",method="get",le="90"} 18
http_request_duration_seconds_bucket{code="200",handler="/metrics",method="get",le="120"} 18
http_request_duration_seconds_bucket{code="200",handler="/metrics",method="get",le="150"} 18
http_request_duration_seconds_bucket{code="200",handler="/metrics",method="get",le="200"} 18
http_request_duration_seconds_bucket{code="200",handler="/metrics",method="get",le="+Inf"} 18
http_request_duration_seconds_sum{code="200",handler="/metrics",method="get"} 0.0480237
http_request_duration_seconds_count{code="200",handler="/metrics",method="get"} 18
```



@bwplotka  
@kakkoyun

## Pitfall #6: Histogram Cardinality Explosion

## Pitfall #7: Poorly Chosen Histogram Buckets

```
# HELP http_request_duration_seconds Tracks the latencies for HTTP requests.
# TYPE http_request_duration_seconds histogram
http_request_duration_seconds_bucket{code="200",handler="/metrics",method="get",le="0.001"} 3
http_request_duration_seconds_bucket{code="200",handler="/metrics",method="get",le="0.00120000000000000001"} 3
http_request_duration_seconds_bucket{code="200",handler="/metrics",method="get",le="0.00140000000000000002"} 10
http_request_duration_seconds_bucket{code="200",handler="/metrics",method="get",le="0.00160000000000000003"} 21
http_request_duration_seconds_bucket{code="200",handler="/metrics",method="get",le="0.00180000000000000004"} 25
http_request_duration_seconds_bucket{code="200",handler="/metrics",method="get",le="0.00200000000000000005"} 29
http_request_duration_seconds_bucket{code="200",handler="/metrics",method="get",le="+Inf"} 49
http_request_duration_seconds_sum{code="200",handler="/metrics",method="get"} 0.10939249999999999
http_request_duration_seconds_count{code="200",handler="/metrics",method="get"} 49
```

## Pitfall #8: Not Initialized Metrics



```
m.connFailedTotal.WithLabelValues("unknown").Add(30)
```

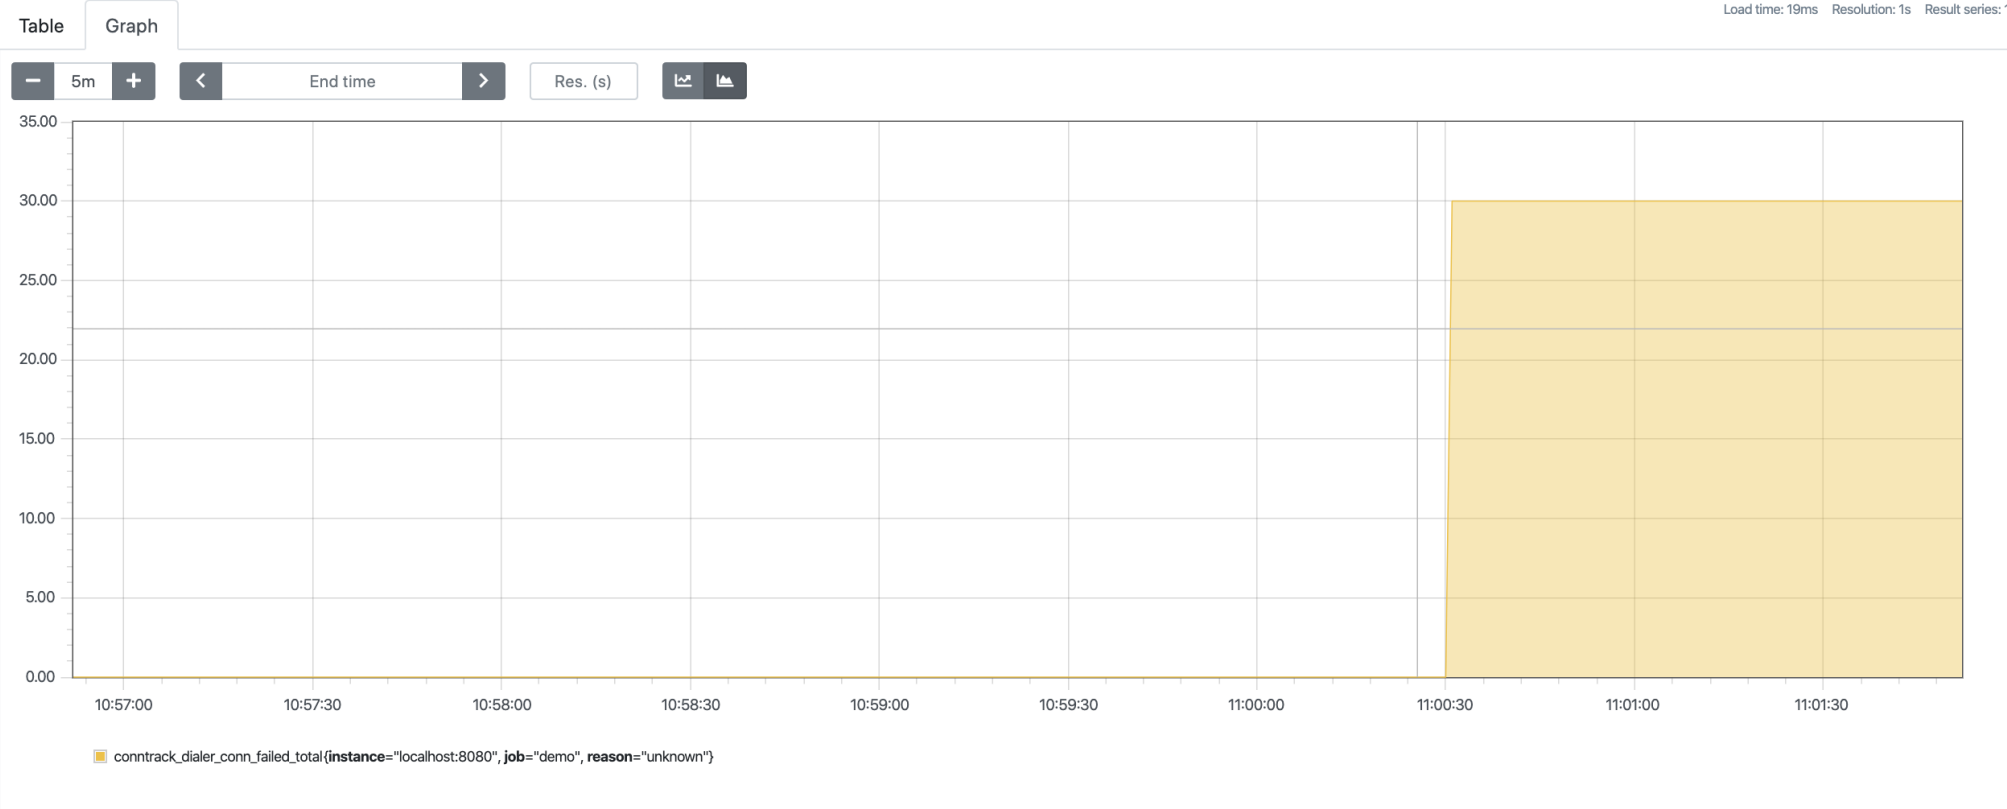
☒ Enable query history

Q

contrack\_dialer\_conn\_failed\_total

G

Execute

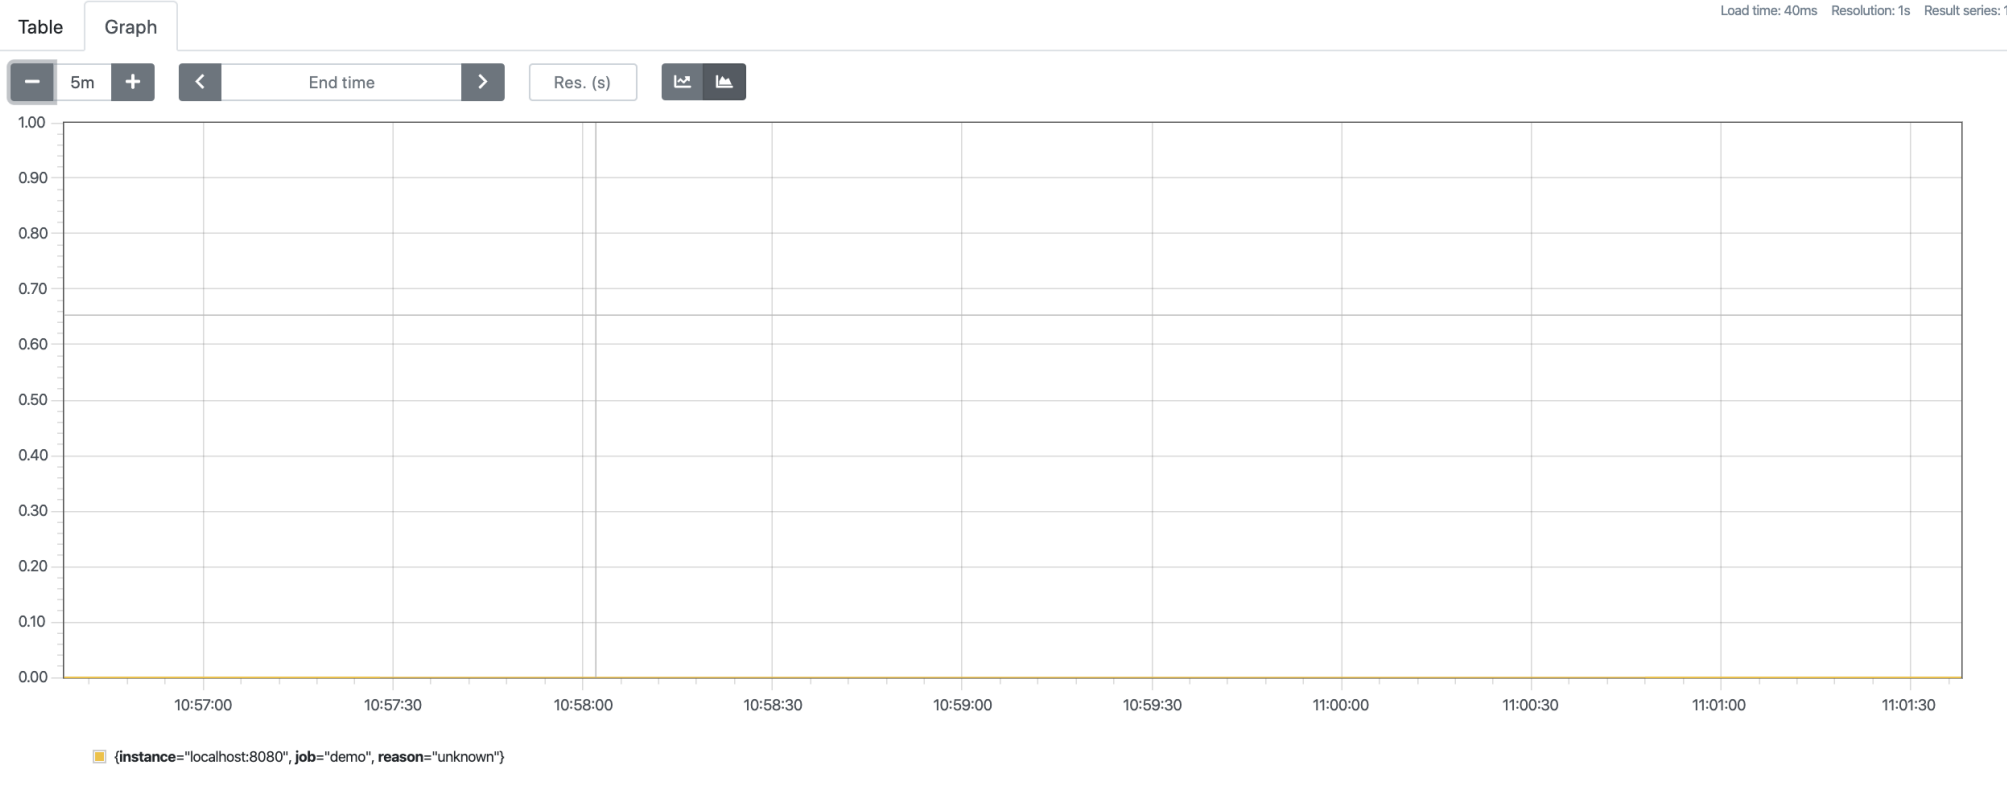
[Remove Panel](#)

☒ Enable query history

Q

increase(contrack\_dialer\_conn\_failed\_total[1m])

Execute



[Remove Panel](#)

## Pitfall #8: Not initializing Metrics

```
m.connFailedTotal.WithLabelValues("unknown").Add(30)
```

☒ Enable query history

Execute

Load time: 54ms Resolution: 3s Result series: 4

Table Graph

15m

+

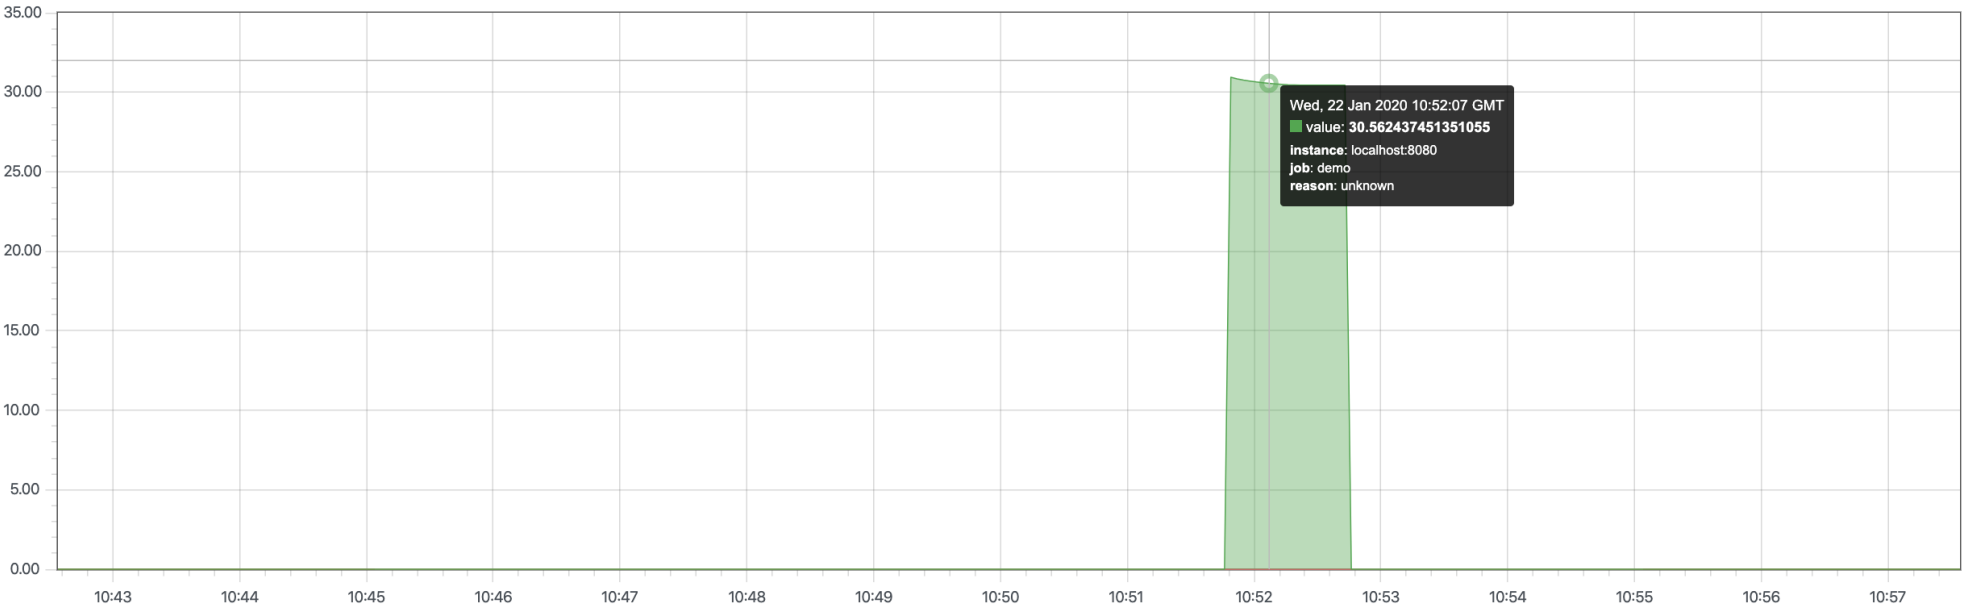
<

2020-01-22 10:57:34

x

>

Res. (s)



- `{instance="localhost:8080", job="demo", reason="refused"}`
- `{instance="localhost:8080", job="demo", reason="resolution"}`
- `{instance="localhost:8080", job="demo", reason="timeout"}`
- `{instance="localhost:8080", job="demo", reason="unknown"}`

## Summary





**Monitoring is essential.**

**Unit Test Your Instrumentation.**

**Avoid Global Registry.**

[Pull requests](#) [Issues](#) [Marketplace](#) [Explore](#)

observatorium / observable-demo

  Unwatch 2  Unstar 4  Fork 1

[Code](#) [Issues 0](#) [Pull requests 0](#) [Actions](#) [Projects 0](#) [Wiki](#) [Security](#) [Insights](#) [Settings](#)

Demo of the instrumented L7 HTTP loadbalancer in Go following best practices

Edit

[Manage topics](#)

12 commits

4 branches

0 packages



0 releases

2 contributors

Apache-2.0

Branch: master New pull request

Create new file Upload files Find file Clone or download

  Add diagram to README (#7)

Latest commit fadac9a 40 minutes ago

cmd/loadbalancer	Added middlewares; fixes; added makefile demo command. (#3)	2 days ago
pkg	Add docs. (#5)	2 days ago
.drone.yml	Clean TODOs (#2)	3 days ago
.errcheck_excludes	Clean TODOs (#2)	3 days ago
.gitignore	Clean TODOs (#2)	3 days ago
.golanci.yml	Clean TODOs (#2)	3 days ago
Dockerfile	Add multistage container builds (#6)	23 hours ago
LICENSE	Initial commit	12 days ago
Makefile	Added middlewares; fixes; added makefile demo command. (#3)	2 days ago
README.md	Add diagram to README (#7)	40 minutes ago
demo-prometheus.yml	Added middlewares; fixes; added makefile demo command. (#3)	2 days ago
go.mod	Added test. (#4)	2 days ago
go.sum	Added middlewares; fixes; added makefile demo command. (#3)	2 days ago

README.md

# observable-demo

semver-0.0.0

maintained yes

build failure

godoc

reference

go report A+

This repository includes simple L7 round robin loadbalancer implementation, instrumented with metrics.

WARNING

: This is not meant to be production ready L7 loadbalancer. It's missing things like proper retrying, DNS discovery, logging, tracing etc, that might be added in future (:

@bwplotka  
@kakkoyun

**Thank you!**



**Red Hat**

<https://github.com/kakkoyun/are-you-testing-your-observability>

@bwplotka  
@kakkoyun



# Reference:

- [Thanos](#)
- [Prometheus](#)
- [Prometheus - client\\_go](#)
- [Prometheus - Histogram](#)
- [Prometheus Histograms – Past, Present, and Future](#)
- [Roboust Perception Blog](#)
- [Why\\_globals are magic](#)
- [Red Method](#)
- [Promcon 2019 - OpenMetrics](#)
- [Gopherize me](#)